



# Smalltalk and Databases

R. Mark Volkmann

Object Computing, Inc.



<https://objectcomputing.com>



[mark@objectcomputing.com](mailto:mark@objectcomputing.com)



[@mark\\_volkmann](https://twitter.com/mark_volkmann)



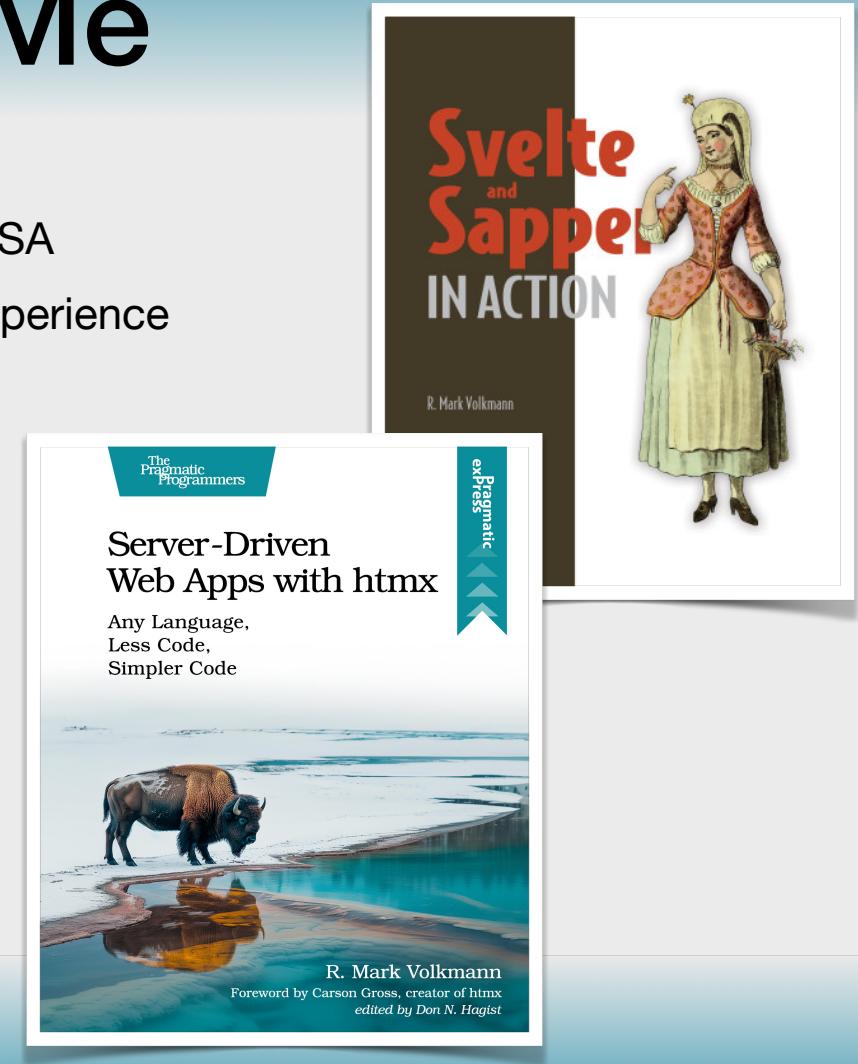
**OBJECT COMPUTING**  
YOUR OUTCOMES ENGINEERED

Slides at <https://github.com/mvolkmann/talks/>

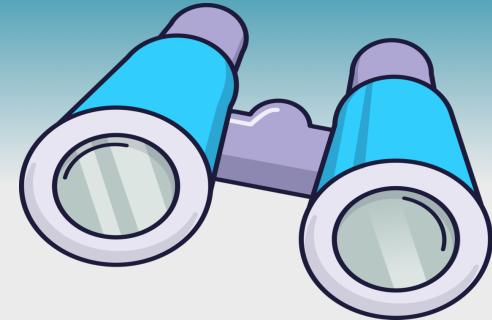


# About Me

- Partner and Distinguished Software Engineer at Object Computing, Inc. in St. Louis, Missouri USA
- 44 years of professional software development experience
- Writer and speaker
- Blog at <https://mvolkmann.github.io/blog/>
- Author of Manning book “Svelte ... in Action”
- Author of Pragmatic Bookshelf book “Server-Driven Web Apps with htmx”



# ODBC Overview



- Acronym for Open DataBase Connectivity
- Defines an API for interacting with databases
  - primarily used with relational databases
- Originally created by Microsoft and Simba Technologies
  - version 1.0 was released in 1992
  - version 4.0 was released in 2016 (latest)
- Supported by a large number of databases
- Somewhat slower than directly accessing a specific kind of database

# Installing ODBC



- In macOS
  - install Homebrew
  - open a Terminal
  - enter `brew install unixodbc`
- This installs
  - `libodbc` library
  - `isql` command
  - `odbcinst` command

**Apologies for Mac-only instructions!**

I need to get access to  
Linux and Windows computers  
to verify their installation steps.

# Installing DB-specific Drivers

- **For Postgres** in macOS
  - in Terminal, enter `brew install pqodbc`
  - creates files `psqlodbcw.so` and `psqlodbca.so` in `/opt/homebrew/lib`
  - “w” stands for “wide”, meaning it supports Unicode characters
  - “a” stands for “ASCII”, meaning it only supports 8-bit characters
- **For MySQL** in macOS
  - in Terminal, enter `brew install mariadb-connector-odbc`
  - creates file `/opt/homebrew/lib/mariadb/libmaodbc.dylib`
- **For SQLite** in macOS
  - in Terminal, enter `brew install sqliteodbc`
  - creates file `/opt/homebrew/lib/libsqlite3odbc.so`



compatible with MySQL

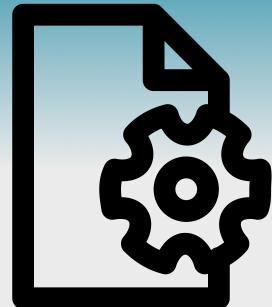
I was unable to get the official MySQL ODBC driver from <https://dev.mysql.com/downloads/connector/odbc/> to work.

**WARNING:** Driver libraries may be installed in `/usr/local/lib` instead of `/opt/homebrew`.

# Defining DSNs

- Acronym for Data Source Name
- DSNs and the drivers they use are configured in text files
- To determine files that must be created
  - in terminal, enter `odbcinst -j`    `-j prints config info`
  - see “DRIVERS” in output
    - mine is `/opt/homebrew/etc/odbcinst.ini`
  - see “USER DATA SOURCES” in output
    - mine is `.odbc.ini` in my home directory
- Create those files with content similar to next slide

# Configuration Files



```
[MySQL]
Description = MySQL ODBC Driver
Driver = /opt/homebrew/lib/mariadb/libmaodbc.dylib

[PostgreSQL]
Description = PostgreSQL ODBC Driver
Driver = /opt/homebrew/lib/psqlodbcw.so

[SQLite]
Description = SQLite ODBC Driver
Driver = /opt/homebrew/lib/libsqLite3odbc.so
```

**WARNING:** Driver libraries may be installed in `/usr/local/lib` instead of `/opt/homebrew`

maps driver names to library files

```
~/.odbc.ini
```

maps data source names to drivers and databases

```
[MySQLTodosDSN]
```

```
Description = MySQL database for a Todo app
Driver = MySQL
Database = todos
User = root
```

alternatively, can set to driver path

```
[PostgresTodosDSN]
```

```
Description = PostgreSQL database for a Todo app
Driver = PostgreSQL
Database = todos
```

```
[SQLiteTodosDSN]
```

```
Description = SQLite database for a Todo app
Driver = SQLite
Database = /Users/volkmannm/todos.db
```

modify for location of your database

# Listing Data Sources

- To list all configured DSNs,  
enter `odbcinst -q -s`
- To list all configured ODBC drivers,  
enter `odbcinst -q -d`
- To view details about a specific data source,  
enter `odbcinst -q -s -n {dsn}`

`-q` for query  
`-s` for source (DSN)  
`-d` for driver  
`-n` for name



# Creating Postgres Database

- Install Postgres: `brew install postgresql@14`
- Create required files: `initdb` directory defaults to /opt/homebrew/var/postgres
- Start database server: `brew services start postgresql@14`
- Create database: `createdb todos`
- Enter interactive mode: `psql -d todos`
- Create todo table
  - ```
create table todos (
    id serial primary key,
    description text unique,
    completed boolean
);
```
- Exit by entering `\q`

# Creating MySQL Database

- Install MySQL commands: `brew install mysql`
- Start database server: `brew services start mysql`
- Enter interactive mode: `/opt/homebrew/bin/mysql -u root`
- Create todo table
  - `create table todos (`  
 `id int auto_increment primary key,`  
 `description varchar(100) not null,`  
 `completed boolean,`  
 `unique (description)`  
`) ;`
  - `must use varchar instead of text`  
to add unique constraint
- Exit by entering `exit`

# Creating SQLite Database

- Install `sqlite3` command: `brew install sqlite`
- cd to directory where database will be created
- Create database and start interactive mode: `sqlite3 todos.db`
- Create todo table
  - `create table todos (`  
    `id integer primary key autoincrement,`  
    `description string unique,`  
    `completed numeric`  
`) ;`
- Exit by pressing `ctrl-d`

# Interactive Mode

- To interactively access a data source,  
enter `isql {dsn}`
  - to list existing tables, enter `help`
  - to see the schema for a specific table (column names and their types),  
enter `help table-name`
  - enter SQL statements like `select * from todos;`
  - to exit, enter `quit` or press ctrl-d

# ODBC from Smalltalk



- In a Workspace, evaluate `Feature require: 'ODBC'`
- Evaluate code like the following

```
| connection resultSet sqlString |  
  
connection := ODBCConnection dsn: #SQLiteTodosDSN user: '' password: ''.  
  
sqlString :=  
    'insert into todos (description, completed) values (''buy milk'', 0)'.  
connection execute: sqlString.  
  
resultSet := connection execute: 'select * from todos'.  
resultSet asTable do: [:row |  
    row at: #description :: print].  
connection close.
```

Only tested in Cuis Smalltalk ...

Will similar code work  
in Squeak and Pharo?

resultSet is  
an instance of  
ODBCResultSet

resultSet asTable returns  
an instance of ODBCResultSet  
which is a subclass of  
OrderedCollection.

# Active Record

- Design pattern described by **Martin Fowler** in book “Patterns of Enterprise Application Architecture”
  - “An object that wraps a row in a database table or view, encapsulates the database access, and adds domain logic on that data.”
  - “The classes match very closely the record structure of an underlying database.”
- Famously implemented as part of **Ruby on Rails** framework
- **Automatically discovers table column names and types**
- Ruby implementation uses `method_missing`
- My Smalltalk implementation uses `doesNotUnderstand:`



# Assumptions

- All tables have an auto-increment `id` column
- All table names are the snake\_case, plural version of the corresponding class name
  - class name -> table name
  - `Todo` -> `todos`
  - `GearShiftKnob` -> `gear_shift_knobs`
  - `Person` -> `people`

“  
Assumptions  
are made and  
most  
assumptions  
are wrong.

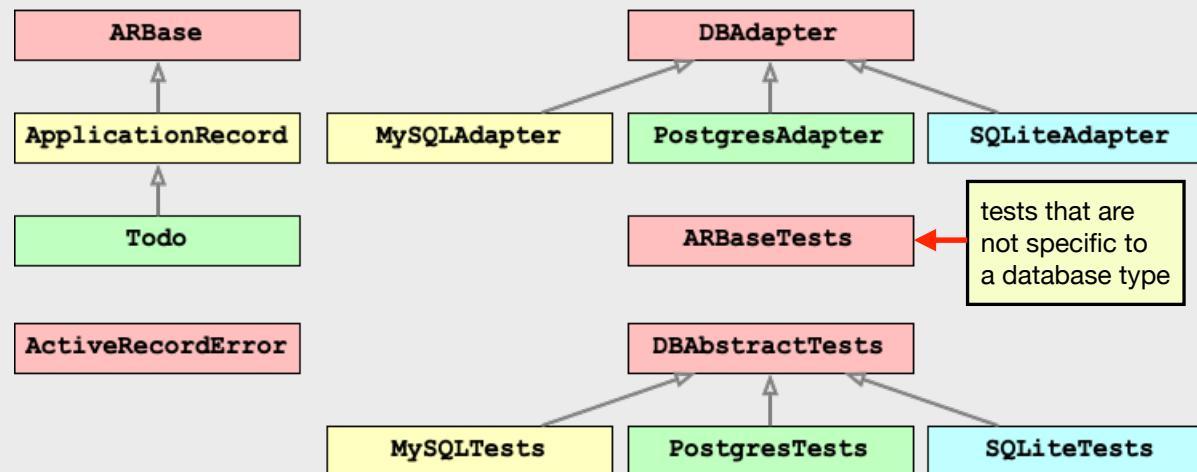
- Albert Einstein

# My Package

- Currently supports MySQL, Postgres, and SQLite
  - to support a new database, implement new subclass of `DBAdapter`
- Currently only tested in Cuis Smalltalk
  - hopefully only minimal changes will be needed for Squeak and Pharo
  - relies on having an ODBC package
- To install in a Cuis image
  - clone repository at  
<https://github.com/mvolkmann/Cuis-Smalltalk-ActiveRecord>
  - open a Workspace and evaluate  
`Feature require: 'ActiveRecord'`



# Classes



# Example Usage ...

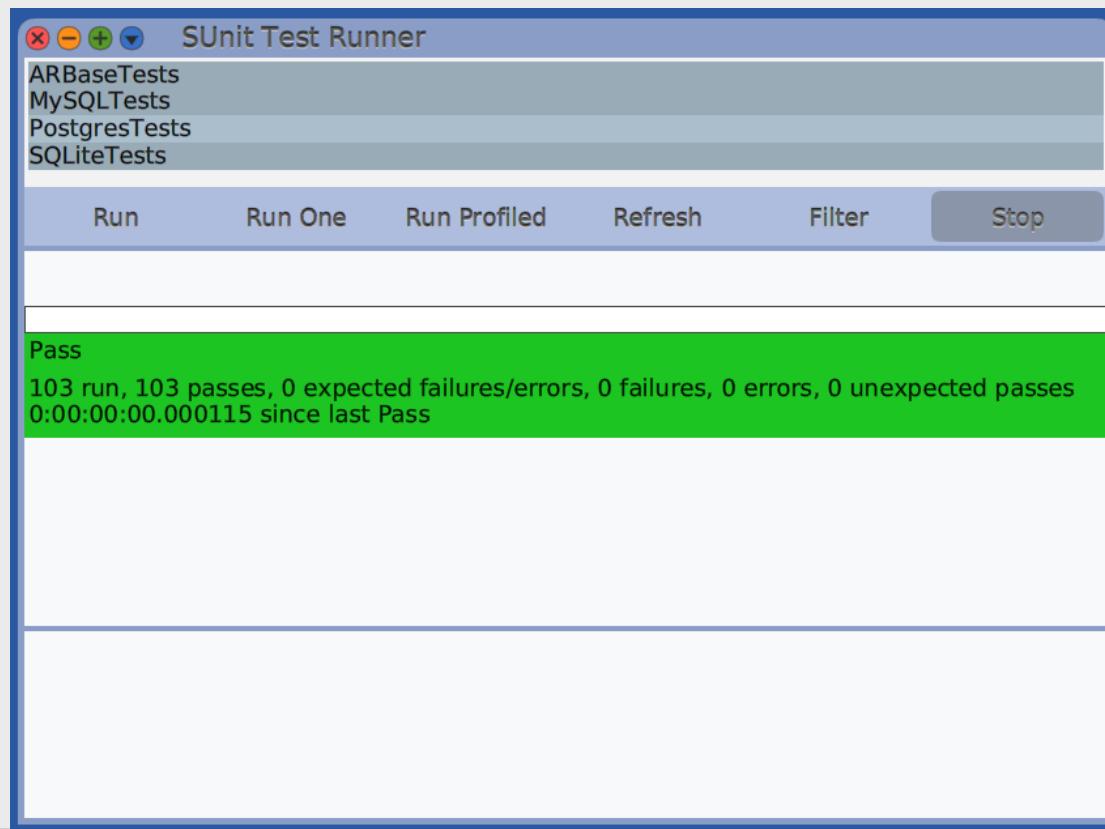
- Mimics Ruby on Rails implementation as closely as possible
- **Todo** class
  - maps to `todos` table
  - only needs to be a subclass of `ApplicationRecord`
  - no instance variables or methods are required

```
ApplicationRecord subclass: #Todo
  instanceVariableNames: ''
  classVariableNames: ''
  poolDictionaries: ''
  category: 'Demo'
```

# ... Example Usage

```
| id todo todos |  
  
ARBase logSQL: true. "logs all SQL statements to Transcript"  
  
Todo establishConnection: #SQLiteTodosDSN dbType: #SQLite.  
id := Todo create: { #description->'eat tacos'. #completed->0 }.  
id := Todo create: { #description->'cut grass'. #completed->1 }.  
id := Todo create: { #description->'buy milk'. #completed->0 }.  
id := Todo create: { #description->'study Smalltalk'. #completed->1 }.  
  
could this be  
determined  
from DSN?  
  
todos := Todo all. "gets all rows"  
todos do: [:todo | todo inspect print].  
outputs lines like  
Todo completed: 1.0, description: 'buy milk', id: 19579  
  
todos := Todo findWhere: {'completed = 0'}. "gets all matching rows"  
  
todo := Todo find: 19. "finds one row by its id"  
todo description: 'go for a run'. "sets description column"  
todo description. "gets description column"  
todo destroy. "deletes row"  
Todo destroyBy: #description->'buy milk'. "deletes all matching rows"  
  
inspect is an  
ActiveRecord method  
that returns a string  
description of a record.
```

# Comprehensive Test Suite



# Rails Features Not Yet Implemented

- Describing table relationships
  - to support joins
- Transactions
  - currently all changes auto-commit
- Overriding assumptions
  - `id` column
  - mapping from classes to table names
- Table creation
- Migrations



# Wrap Up

It's a good start, but  
there is much more work to do  
to match the functionality of the  
Ruby on Rails implementation.